Predicting Real Estate Sales in NYC:

Combining Machine Learning with Spatial Analysis to

Boost Predictive Accuracy

By

Tim Kiely

Thesis Project

Submitted in partial fulfillment of the

requirements for the degree of

MASTER OF SCIENCE IN DATA SCIENCE

Northwestern University

August, 2018

Nathaniel D. Bastian, PhD, First Reader

Candice Bradley, Second Reader

Contents

1 Introduction

Income inequality may be one of the most pressing challenges of our time, yet, its causes remain unclear. What is clear is that, increasingly, the places where people live "over time are becoming more segregated by income, due in part to macro-level increases in income inequality" (Zuk 2015). Income inequality has many well-explored dimensions: social, economic, political, philosophical. In this paper, we focus on an under-explored dimension, one with potential for tremendous impact: geo-spatial. The locations where people live (by choice or not) have a dramatic impact on their quality of life and ability to progress. Yet, many who are displaced by gentrification find themselves in self-reinforcing cycles of displacement and poverty.

Researchers at the Urban Institute (Solomon Greene and Lei 2016) recently identified the socio-economic phenomenon of "Economic Exclusion" as a direct cause of income inequality in the US. "Economic Exclusion" can be defined as follows: vulnerable populations– disproportionately communities of color, immigrants, refugees, and women–who are physically displaced by local economic prosperity can enter into a gradual cycle of diminished access to good jobs, good schools, health care facilities, public spaces and other physically proximate benefits. Diminished access leads to more poverty, which leads to more displacement. Such self-reinforcing poverty gradually exacerbates income inequality over the course years and even generations.

One practical way to combat Economic Exclusion is to focus on preventing displacement, i.e., the physical relocation of populations away from economic resources. As defined by Clay (1979), Displacement is the negative consequence of gentrification. Reliably predicting gentrification would be a valuable tool for preventing displacement at an early stage, however, such a task has proven difficult historically.

When an area experiences economic growth, increased housing demands and subsequent affordability pressures can lead to voluntary or involuntary relocation of low-income families and small businesses. Government agencies and nonprofits tend to intervene once displacement

is already underway, and after-the-fact interventions can be costly and ineffective. As explained by Solomon Greene and Lei (2016), there are several preemptive actions that can be deployed to stem divestment and ensure that existing residents benefit from new investments. Not unlike medical treatment, early detection is the key to success.

Consequently, in 2016, the Urban Institute put forth a call for research into the creation of "neighborhood-level early warning and response systems that can help city leaders and community advocates get ahead of neighborhood changes" (Solomon Greene and Lei 2016). This paper explores a technique to answer that call in part using free, open data and open-source software.

Many government agencies have already become competent at using predictive modeling to identify and address socio-economic challenges at the individual-level, ranging from prescription drug abuse to homelessness to recidivism (Ritter 2013). However, few, if any, examples exist of large-scale, systematic applications of data analysis to aid vulnerable populations experiencing displacement. This paper belongs to an emerging trend known as the "science of cities" which aims to use large data sets and advanced simulation and modeling techniques to understand urban patterns and how cities function (Batty 2013).

In this paper, we explore techniques that can dramatically boost the accuracy of existing gentrification prediction models. We use real estate transactions, both their occurrence (probability of sale) and their dollar amount (sale price per square foot) as a proxy for gentrification. We explain how this predictive technique may be applied to practically combat Economic Exclusion, a precursor and contributor to Income Inequality. The technique marries the use of machine-learning predictive modeling (Random Forrest) with "spatial-lag" features typically seen in geographically-weighted regressions (GWR). We find that, while the addition of many new variables to a modeling data set can inhibit the models' ability to generalize into the future, spatial-lag features 1) consistently outperform zip-code level aggregation features, and 2) consistently outperform all models for specific property types. We conclude that spatial-lag features, while computationally expensive, can be used to greatly increase

the accuracy of spatially-conscious predictive models.

2 Literature Review

The literature review for this paper discusses the concept of Economic Displacement as it has been addressed in academia, primarily in relation to the study of gentrification. We also examine "mass appraisal techniques", which are automated analytical techniques used for valuing large numbers of real estate properties. Finally, we will briefly examine the literary treatment of machine learning as it relates to the problem of predicting gentrification and/or Economic Displacement.

2.1 How Has Economic Displacement Been Addressed in the Past?

Economic Displacement has been intertwined with the study of gentrification since shortly after the latter became academically relevant in the 1960's. The term "gentrification" was first used by Ruth Glass in 1964 to described the "gentry" in low income neighborhoods in London. Gentrification was originally understood as a "tool of revitalization for declining neighborhoods" (Zuk 2015), however, in 1979 Phillip Clay made the distinction between two types of revitalization: "incumbent upgrading" and "gentrification", noting that Economic Displacement was the negative consequence of the latter (Clay 1979). Today, the term has evolved to describe "a spatial organization and re-organization of human dwelling and activity" (Zuk 2015). Specific to cities, gentrification is thought of as "the transformation of a working-class or vacant area of the central city into middle-class residential or commercial use" (Lees 2008).

Studies of gentrification and displacement generally take two approaches in the literature: supply-side and demand-side, or "the flows of capital versus flows of people to neighborhoods", respectively (Zuk 2015). Supply side arguments for gentrification tend to focus on "private

capital investment, public policies, and public investments" (Zuk 2015), and are much more often the subject of academic literature on Economic Displacement. This kind of research is more common because it has the advantage of being more directly linked to influencing public policy (as opposed to controlling the flows of people). According to Dreier (2004), public policies that have been negatively linked to Economic Displacement have been, among others, automobile-oriented transportation infrastructure spending and mortgage interest tax deductions for home owners. Others that have argued for supply-side gentification include Smith (1979), who stated that the return of capital from the suburbs to the city, or the "political economy of capital flows into urban areas" are what primarily drive both the positive and negative consequences of urban gentrification.

More recently, income inequality has been explored as a major consequence of Economic Displacement. Specifically, "higher compensation in the top quintile and the lack of jobs for the bottom quintile" (Reardon 2011); (Watson 2009). The concentration of wealth allows "certain households to sort themselves according to their preferences – and control local political processes that continue exclusion" (Reardon 2011). This results in a self-reinforcing feedback loop where wealthier households influence public policy toward their self interest. Gentrification prediction tools could be used to help break such feedback loops through early identification and intervention.

Many studies conclude that gentrification in most forms leads to exclusionary economic displacement, however, Zuk (2015) characterizes the results of many recent studies as "mixed, due in part to methodological shortcomings". In this paper, we attempt to further the understanding of gentrification prediction by demonstrating a technique to better predict real estate sales in New York City.

2.2   A Review of Mass Appraisal Techniques

Much of the research on predicting real estate prices has been in service of creating mass appraisal models. Mass appraisal models are most commonly used by local governments

for the purpose of collecting taxes from property owners. Mass appraisal models share many characteristics with predictive machine learning models, in that they are data-driven, standardized methods that employ statistical testing (Eckert 1990). A variation on mass appraisal models are the "automated valuation models" (AVM), which use "often the same methodological framework of mass appraisal... a statistical model and a large amount of property data to estimate the market value of an individual property or portfolio of properties" (d'Amato 2017).

Scientific mass appraisal models date back to 1936 with the reappraisal of St. Paul, Minnesota (Silverherz 1936). Since that time, and accelerated with the advent of computers, much statistical research has been done relating property values and rent prices to various characteristics of those properties, including characteristics of their surrounding area. Multiple regression analysis (MRA) has been the most common set of statistical tools used in mass appraisal, including Maximum Likelihood, Weighted Least Squares, and the most popular, Ordinary Least Squares (OLS) (d'Amato 2017). The primary drawbacks of MRA techniques are "excessive multicollinearity among attributes" and "spatial autocorrelation among residuals" (d'Amato 2017). Another group of models that seek to correct for spatial dependence are known as Spatial Auto Regressive models (SAR), chief among them the Spatial Lag Model, which aggregates weighted summaries of nearby properties in order to create independent regression variables (d'Amato 2017).

So-called "Hedonic" regression models seek to decompose the price of a good based on the intrinsic and extrinsic components. Koschinsky (2012) is a recent and thorough discussion of parametric hedonic regression techniques. Some of the variables included in Koschinsky's models are derived from nearby properties, similar to the technique used in this paper, and these variables were found to be predictive. The real estate hedonic model as defined by Koschinsky describes the price of a property as:

$$P_i = P(S_i, N_i, L_i)$$

Where P_i represents the price of house i, which is a composite good comprised of a vector of structural characteristics S, a vector of social and neighborhood characteristics N, and a vector of locational characteristics L. Specifically, the model calculates spatial lags on properties of interest using neighboring properties within 1,000 feet of a sale. The derived variables include characteristics like average age, quantity of poor condition homes, percent of homes with electric heating, construction grade, etc., within 1,000 feet of the property in question. Koschinsky found that in all cases, "the relation between a home's price and the average price of its neighboring homes is characterized by positive spatial autocorrelation" meaning that homes near each other were typically similar to each other and priced accordingly. Koschinsky concluded that locational characteristics should be valued at least as much "if not more" than important structural characteristics.

As recently as 2015, much research has dealt with mitigating the drawbacks of MRA, including the use of multi-level hierarchical models. Fotheringham (2015) explored the combination of Geographically Weighted Regression (GWR) with time-series forecasting to predict home prices over time. GWR is a variation on OLS that allows for "adaptive bandwidths" of local data to be included, i.e., for each estimate, the number of data points included varies and can be optimized using cross-validation.

2.3 A Review of Machine Learning Applied to the Prediction of Gentrification

Both Mass Appraisal techniques and Automated Valuation Modeling seek to predict real estate prices using data and statistical methods, however, traditional techniques typically fall short. This is because property valuation is inherently a "chaotic" process that does not lend itself to binary or linear analysis (Zuk 2015). The value of any given property is a complex combination of perceived value and speculation. The value of any building or plot of land belongs to a rich network where decisions about and perceptions of neighboring properties influence the final market value. Guan et al. (2014) compared traditional MRA techniques to

alternative "data mining techniques" resulting in "mixed results". However, as Helbich (2013) states, hedonic pricing models "can be improved in two ways: (a) Through novel estimation techniques, and (b) by ancillary structural, locational, and neighborhood variables on the basis of Geographic Information System (GIS)". Recent research generally falls into these two buckets: better analysis algorithms and/or better data.

In the "better data" category, researchers have been striving to introduce new independent variables to increase the accuracy of predictive models. Dietzell (2014) successfully used internet search query data provided by Google Trends to serve as a sentiment indicator and improve commercial real estate forecasting models. Pivo and Fisher (2011) examined the effects of walkability on property values and investment returns. Pivo found that on a 100-point scale, a 10-point increase in walkability increased property investment values by up to 9%.

Research into better prediction algorithms do not necessarily happen at the exclusion of "better data". For example, Fu (2014) created a prediction algorithm, called "ClusRanking", for real estate in Beijing, China. ClusRanking first estimates neighborhood characteristics using taxi cab traffic vector data, specifically as they relate to accessibility to "business areas". Then, the algorithm performs a rank-ordered prediction of investment returns segmented into five categories. Similar to Koschinsky (2012), though less formally stated, Fu (2014) thought of a property's value as a composite of individual, peer and zone characteristics. In the predictive model, Fu includes characteristics of the neighborhood (individual), the values of its nearby properties (peer), and the prosperity of the affiliated latent business area (zone) based on taxi cab data (Fu 2014).

Several other recent studies compare various "advanced" statistical techniques and algorithms either to other advanced techniques or to traditional ones. Most studies conclude that the advanced, non-parametric techniques outperform traditional parametric techniques, while several conclude that the Random Forest algorithm is particularly well-suited to predicting real estate values.

Kontrimasa (2011) compares the accuracy of linear regression against the SVM technique and found the latter to outperform. Schernthanner H. (2016) compared traditional linear regression techniques to several techniques such as krigging (stochastic interpolation) and Random Forest. They concluded that the more advanced techniques, particularly random forest, are sound and more accurate when compared to traditional statistical methods. Antipov and Pokryshevskaya (2012) came to a similar conclusion about the superiority of Random Forest for real estate valuation after comparing 10 algorithms: multiple regression, CHAID, Exhaustive CHAID, CART, 2 types of k-Nearest Neighbors, Multilayer Perceptron artificaial neural network (ANN), Radial Basis Function neural network (RBF), Boosted Trees and finally Random Forest.

Guan et al. (2014) compared three different approaches to defining spatial neighbors: a simple radius technique, a k-nearest neighbors technique using only distance and a k-nearest neighbors technique using all attributes. Interestingly, the location-only KNN models performed best, although by a slight margin. Park (2015) developed several housing price prediction models based on machine learning algorithms including C4.5, RIPPER, Naive Bayesian, and AdaBoost. By comparing the models' classification accuracy performance, the experiments demonstrate that the RIPPER algorithm, based on accuracy, consistently outperformed the other models in the performance of housing price prediction. Rafiei (2016) employed a restricted boltzmann machine (neural network with back propagation) to predict the sale price of residential condos in Tehran, Iran, using a non-mating genetic algorithm for dimensionality reduction with a focus on computational efficiency. The paper concludes that two primary strategies help in this regard: weighting property sales by temporal proximity (sales which happened closer in time are more alike), and also using a learner to accelerate the recognition of important features. The paper compares this technique to several other common neural network approaches and finds that while not necessarily the only way to get the best answer, it is the fastest way to get to the best answer.

Finally, it should be noted that many studies, whether exploring advanced techniques,

new data, or both, rely on aggregation of data by some arbitrary boundary. For example, Turner and Snow (2001) predicted gentrification in the Washington, D.C. metro area by ranking census tracts in terms of development. Chapple (2009) created a gentrification "early warning system" by identifying low income census tracts in central city locations. Barry Bluestone & Chase Billingham (2010) analyzed 42 census block groups near rail stations in 12 metro areas across the United States, studying changes between 1990 and 2000 for neighborhood socioeconomic and housing characteristics. All of these studies, and many more, relied on aggregation of data at the census-tract or census-block level. In contrast, this paper compares boundary-aggregation techniques (specifically, aggregating by zip codes) to spatial-lag techniques and finds the spatial lag techniques to generally outperform.

3   Data and Methodology

3.1   Methodology Overview

Our goal is to compare "spatially-concious" machine learning predictive models to traditional feature engineering techniques. To accomplish this comparrisson, we create three separate modeling data sets:

- **Base modeling data:** includes building characteristics such as size, taxable value, usage and others
- **Zip Code modeling data:** includes base data as well as aggregations of data at the zip-code level
- **Spatial Lag modeling data:** includes base data as well as aggregations of data within 500-meters of each building

The second and third modeling datasets are incremental variations of the first, using competing feature engineering techniques to extract additional predictive power from the data. To accomplish this, we combine three open-source data repositories provided by New York City via nyc.gov and data.cityofnewyork.us. Our base modeling data set includes all building

Table 3.1: Six Predictive Models

#	Model	Model Task	Data	Outcome Var	Outcome Type	Eval Metric
1	Probability of Sale	Classification	Base	Building Sold	Binary	AUC
2	Probability of Sale	Classification	Zip Code	Building Sold	Binary	AUC
3	Probability of Sale	Classification	Spatial Lags	Building Sold	Binary	AUC
4	Sale Price	Regression	Base	Sale Price per SF	Continuous	RMSE
5	Sale Price	Regression	Zip Code	Sale Price per SF	Continuous	RMSE
6	Sale Price	Regression	Spatial Lag	Sale Price per SF	Continuous	RMSE

records and associated sales information from 2003-2017. For each of the three modeling dataset, we also compare 2 predictive modeling tasks, using a different outcome variable for each:

1) **Classification task: Probability of Sale** The probability that a given property will sell in a given year

2) **Regression task: Sale Price per square foot** Given that a property sells, how much is the sale price per square foot?

The six distinct modeling tasks/data combinations are shown in table 3.1.

We conduct our analysis in a two-stage process. In Stage 1, we use the Random Forrest algorithm to evaluate the data's suitability for our feature engineering assumptions. In Stage 2, we create subsets of the modeling datasets based on analysis conducted in Stage 1. We then compare the performance of different algorithms across all modeling datasets and prediction tasks. The following is an outline of our complete analysis process:

Stage 1: Random Forrest, all data

1) Create a "base" modeling dataset by sourcing and combining building characteristic and sales data from open-source New York City repositories

2) Create a "zip" modeling dataset by aggregating the base data at a Zip-Code level and appending these features to the base data

3) Create a "spatial lag" modeling dataset by aggregating the base data within 500 meters of each building and appending these features to the base data

4) Train a Random Forrest model on all three datasets, for both classification and regression tasks

5) Evalutate the performance of the various Radom Forrest models on hold-out test data

6) Analyze the prediction results by building type and geography, identifying those buildings for which our feature-engineering assumptions (e.g., 500 meter radii spatial lags) are most appropriate

Stage 2: Many algorithms, refined data

7) Create subsets of the modeling data based on analysis conducted in Stage 1

8) Train machine learning models on the refined modeling data sets using several algorithms, for both classification and regression tasks

9) Evaluate the performance of the various models on hold-out test data

10) Analyze the prediction results of the various algorithm/data/task combinations

3.2 Data

3.2.1 Data Sources

The New York City government makes available an annual data set which describes all tax lots in the five boroughs. The Primary Land Use and Tax Lot Output data set, known as PLUTO[1], contains a single record for every tax lot in the city along with a number of building and tax-related attributes such as Year Built, Assessed Value, Square Footage, number of stories, and many more. At the time of this writing, NYC has made this data set available for all years between 2002-2017, excluding 2008. For convenience, we also exclude the 2002 data set from our analysis because corresponding sales information is not available for that year. Importantly for our analysis, the latitude and longitude of the tax lots are also made available, allowing us to locate in space each building and to build geospatial features from the data.

---

[1]https://www1.nyc.gov/site/planning/data-maps/open-data/bytes-archive.page?sorts%5Byear%5D=0

Ultimately, we are interested in both the ocuurance and the amount of real estate sales transactions. Sales transactions are also made available by the New York City government, known as NYC Rolling Sales Data[2]. At the time of this writing, sales transactions are available for the years 2003-2017. The sales transactions data contains additional data fields describing time, place, and amount of sale as well as additional building characteristics. Crucially, the sales transaction data does not include geographical coordinates, making it impossible to perform geospatial analysis without first mapping the sales data to PLUTO.

Prior to mapping to PLUTO, the sales data must first be transformed to include the proper mapping key. New York City uses a standard key of Borough-Block-Lot to identify tax lots in the data. For example, 31 West 27th Street is located in Manhattan, on block 829 and lot 16, therefore, its Borough-Block-Lot (BBL) is 1_829_16 (the 1 represents Manhattan). The sales data contains BBL's at the building level, however, the sales transactions data does not appropriately designate condos as their own BBL's. Mapping the sales data directly to the PLUTO data results in a mapping error rate of 23.1%. Therefore, the sales transactions data must first be mapped to another data source, the NYC Property Address Directory, or PAD[3], which contains an exhaustive list of all BBL's in NYC. Once the sales data is combined with PAD, the data can be mapped to PLUTO with an error rate of 0.291% (See: Figure 3.1).

---

[2]http://www1.nyc.gov/site/finance/taxes/property-annualized-sales-update.page
[3]https://data.cityofnewyork.us/City-Government/Property-Address-Directory/bc8t-ecyu/data

	PLUTO	PAD	Rolling Sales Data	Mapping Error Rate

Figure 3.1: Overview of Data Sources

After the Sales Transactions data has been mapped to PAD, it can then be mapped to PLUTO. The sales data is normalized and filtered so that only BBL's with less than or equal to 1 transactions in a year occur. The final data set is an exhaustive list of all tax lots in NYC for every year between 2003-2017, whether that building was sold, for what amount, and several other additional variables. A description of all variables can be seen in Table 3.2.

3.2.2 Global filtering of the data

We will only include building categories of significant interest in the modeling data. Generally speaking, by significant interest we are referring to building types that are regularly bought and sold on the free market. These include residences, office buildings and industrial buildings, and exclude things like government-owned buildings and hospitals. We also exclude hotels as they tend to be comparatively rare in the data and exhibit unique sales characteristics. The included building types are displayed in Table 3.3.

The data is further filtered to include only records with equal to or less than 2 buildings per tax lot. This effectively excludes large outliers in the data, including multi-building tax lots such as the World Trade Center and Stuyvesant Town. The global filtering of the data set reduces the base modeling data from 12,012,780 records down to 8,247,499, retaining 68.6%% of the original data.

Table 3.2: Description of Base Data

variable	type	nobs	mean	sd	mode	min	max	median	n_missing
Annual_Sales	Numeric	12,012,780	1.60	8.24	NA	1.00	2,591.00	1.00	11,208,593
AssessLand	Numeric	12,012,780	93,492.94	2,870,654.30	103050	0.00	2,146,387,500.00	10,348.00	65
AssessTot	Numeric	12,012,780	302,375.19	4,816,339.21	581400	0.00	2,146,387,500.00	25,159.00	1,703,150
BldgArea	Numeric	12,012,780	6,228.48	70,160.95	18965	0.00	49,547,830.00	2,050.00	45
BldgDepth	Numeric	12,012,780	45.54	34.42	50	0.00	9,388.00	42.00	44
BldgFront	Numeric	12,012,780	25.38	32.51	100	0.00	9,702.00	20.00	44
Block	Numeric	12,012,780	5,297.42	3,694.56	1	0.00	71,724.00	4,799.00	44
BoroCode	Numeric	12,012,780	3.46	1.02	5	1.00	5.00	4.00	47
BsmtCode	Numeric	12,012,780	2.37	1.98	0	0.00	3,213.00	2.00	859,406
BuiltFAR	Numeric	12,012,780	1.11	9.94	3.32	0.00	8,695.00	0.76	850,554
ComArea	Numeric	12,012,780	2,159.69	58,192.19	18965	0.00	27,600,000.00	0.00	44
CommFAR	Numeric	12,012,780	0.21	1.10	3.4	0.00	15.00	0.00	7,716,603
CondoNo	Numeric	12,012,780	8.13	125.78	0	0.00	30,000.00	0.00	1,703,113
Easements	Numeric	12,012,780	0.01	2.18	0	0.00	7,500.00	0.00	48
ExemptLand	Numeric	12,012,780	37,073.07	2,718,193.73	0	0.00	2,146,387,500.00	1,290.00	65
ExemptTot	Numeric	12,012,780	107,941.41	3,522,172.08	0	0.00	2,146,387,500.00	1,360.00	1,703,149
FacilFAR	Numeric	12,012,780	2.23	1.69	4.8	0.00	15.00	2.00	7,716,603
FactryArea	Numeric	12,012,780	126.48	3,889.92	0	0.00	1,324,592.00	0.00	850,555
GarageArea	Numeric	12,012,780	130.46	5,154.00	0	0.00	2,677,430.00	0.00	850,554
GROSS SQUARE FEET	Numeric	12,012,780	4,422.72	45,691.20	NA	0.00	14,962,152.00	1,920.00	11,217,669
lat	Numeric	12,012,780	40.69	0.08	40.6386175499986	40.11	40.91	40.69	427,076
lon	Numeric	12,012,780	-73.92	0.12	-74.0754964873625	-77.52	-73.70	-73.91	427,076
Lot	Numeric	12,012,780	114.74	655.29	10	0.00	9,999.00	38.00	44
LotArea	Numeric	12,012,780	7,852.07	362,618.31	5716	0.00	214,755,710.00	2,514.00	44
LotDepth	Numeric	12,012,780	104.01	68.77	84	0.00	9,999.00	100.00	45
LotFront	Numeric	12,012,780	39.98	73.95	112.58	0.00	9,999.00	25.00	44
LotType	Numeric	12,012,780	4.72	0.78	5	0.00	9.00	5.00	865,340
NumBldgs	Numeric	12,012,780	1.17	3.87	1	0.00	2,740.00	1.00	46
NumFloors	Numeric	12,012,780	2.28	1.90	4	0.00	300.00	2.00	44
OfficeArea	Numeric	12,012,780	741.81	21,566.05	0	0.00	5,009,319.00	0.00	850,556
OtherArea	Numeric	12,012,780	672.95	49,848.48	0	0.00	27,600,000.00	0.00	850,555
ProxCode	Numeric	12,012,780	1.49	1.90	1	0.00	5,469.00	1.00	197,927
ResArea	Numeric	12,012,780	3,921.11	31,881.56	0	0.00	35,485,021.00	1,776.00	44
ResidFAR	Numeric	12,012,780	1.37	1.38	2.43	0.00	12.00	0.90	7,716,603
RetailArea	Numeric	12,012,780	308.82	14,393.95	6965	0.00	21,999,988.00	0.00	850,554
SALE PRICE	Numeric	12,012,780	884,035.89	13,757,705.99	NA	0.00	4,111,111,766.00	319,000.00	11,208,593
sale_psf	Numeric	12,012,780	219.72	5,153.01	NA	0.00	1,497,500.00	114.13	11,250,396
SALE_YEAR	Numeric	12,012,780	2,009.37	4.66	NA	2,003.00	2,017.00	2,009.00	11,208,593
Sold	Numeric	12,012,780	0.07	0.25	0	0.00	1.00	0.00	0
StrgeArea	Numeric	12,012,780	168.96	5,810.14	12000	0.00	1,835,150.00	0.00	850,554
TOTAL_SALES	Numeric	12,012,780	884,035.89	13,757,705.99	NA	0.00	4,111,111,766.00	319,000.00	11,208,593
UnitsRes	Numeric	12,012,780	3.96	36.44	0	0.00	20,811.00	1.00	45
UnitsTotal	Numeric	12,012,780	4.32	41.84	1	0.00	44,276.00	2.00	47
Year	Numeric	12,012,780	2,010.15	4.43	2017	2,003.00	2,017.00	2,011.00	0
YearAlter1	Numeric	12,012,780	158.51	539.53	2000	0.00	2,017.00	0.00	45
YearAlter2	Numeric	12,012,780	20.49	201.53	0	0.00	2,017.00	0.00	48
YearBuilt	Numeric	12,012,780	1,829.84	448.73	1884	0.00	2,040.00	1,930.00	47
ZipCode	Numeric	12,012,780	11,006.64	537.34	10301	0.00	11,697.00	11,221.00	59,956
Address	Character	12,012,780	NA	NA	139 BAY STREET	NA	NA	NA	17,902
AssessTotal	Character	12,012,780	NA	NA	NA	NA	NA	NA	10,309,712
bbl	Character	12,012,780	NA	NA	5_1_10	NA	NA	NA	0
BldgClass	Character	12,012,780	NA	NA	E1	NA	NA	NA	16,372
Borough	Character	12,012,780	NA	NA	SI	NA	NA	NA	0
BUILDING CLASS AT PRESENT	Character	12,012,780	NA	NA	NA	NA	NA	NA	11,219,514
BUILDING CLASS AT TIME OF SALE	Character	12,012,780	NA	NA	NA	NA	NA	NA	11,208,593
BUILDING CLASS CATEGORY	Character	12,012,780	NA	NA	NA	NA	NA	NA	11,208,765
Building_Type	Character	12,012,780	NA	NA	E	NA	NA	NA	16,372
CornerLot	Character	12,012,780	NA	NA	NA	NA	NA	NA	11,163,751
ExemptTotal	Character	12,012,780	NA	NA	NA	NA	NA	NA	10,309,712
FAR	Character	12,012,780	NA	NA	NA	NA	NA	NA	11,162,270
IrrLotCode	Character	12,012,780	NA	NA	Y	NA	NA	NA	16,310
MaxAllwFAR	Character	12,012,780	NA	NA	NA	NA	NA	NA	4,296,221
OwnerName	Character	12,012,780	NA	NA	139 BAY POINTE PROPER	NA	NA	NA	137,048
OwnerType	Character	12,012,780	NA	NA	NA	NA	NA	NA	10,445,328
TAX CLASS AT PRESENT	Character	12,012,780	NA	NA	NA	NA	NA	NA	11,219,514
TAX CLASS AT TIME OF SALE	Character	12,012,780	NA	NA	NA	NA	NA	NA	11,208,593
ZoneDist1	Character	12,012,780	NA	NA	C4-2	NA	NA	NA	18,970
ZoneDist2	Character	12,012,780	NA	NA	NA	NA	NA	NA	11,715,653

Table 3.3:   Included Building Cateogory Codes

Category	Description
A	ONE FAMILY DWELLINGS
B	TWO FAMILY DWELLINGS
C	WALK UP APARTMENTS
D	ELEVATOR APARTMENTS
F	FACTORY AND INDUSTRIAL BUILDINGS
G	GARAGES AND GASOLINE STATIONS
L	LOFT BUILDINGS
O	OFFICES

3.2.3   Exploratory Data Analysis

The data contain building and sale records across the five boroughs of New York City for the years 2003-2017. One challenge with creating a predictive model of real estate sales data is the heterogeneity within the data in terms of frequency of sales and sale price. These two metrics (sale occurance and amount) vary greatly across year, borough and building classes (among other attributes). Table 3.4 displays statistics which describe the base data set (pre-filtered) by year. Note how the frequency of transactions (# of Sales) and the sale amount (Median Sale \$/SF) tend to covary, particularly through the downturn of 2009-2012. This may be due to the fact that the reltive size of transactions tends to decrease as capital becomes more constrained.

Similar variance can be seen across asset types. Table 3.5 shows all buildings classes in the 2003-2017 period. Unsurprisingly, residences tend to have the highest volume of sales while offices tend to have the highest sale prices.

Sale price per square foot, in particular, varies greatly across geography and assett class. Table 3.6 shows the breakdown of sales prices by borough and asset class. Manhattan tends to command the highest sale-price-per-square foot across assett types. "Comercial" asset types such as Office and Elevator Apartments tend to fetch much lower price-per-square foot than do residential classes such as one and two-family dwellings. Table 3.7 shows the number

Table 3.4:   Sales By Year

Year	N	# Sales	Median Sale	Median Sale $/SF
2003	850515	78919	$218,000	$79.37
2004	852563	81794	$292,000	$124.05
2005	854862	77815	$360,500	$157.76
2006	857473	70928	$400,000	$168.07
2007	860480	61880	$385,000	$139.05
2009	860519	43304	$245,000	$41.25
2010	860541	41826	$273,000	$75.35
2011	860320	40852	$263,333	$56.99
2012	859329	47036	$270,708	$52.72
2013	859372	50408	$315,000	$89.44
2014	858914	51386	$350,000	$115.71
2015	859464	53208	$375,000	$135.62
2016	859205	53772	$385,530	$147.06
2017	859223	51059	$430,000	$171.71

of transactions across the same dimensions.

3.3   Feature Engineering

3.3.1   Base Modeling Data

The base modeling data set is constructed by combining several open-source data repositories, outlined in the Data Sources section. In addition to the data provided by New York City, several additional features are engineered and appended. A summary table of the additional features are presented in Table 3.8. A binary variable is created to indicate whether a tax lot has a building on it (i.e., whether it is an empty plot of land or not). In addition, building types are quantified by what percent of the square footage belongs to the major property types: Commercial, Residential, Office, Retail, Garage, Storage, Factory and Other.

Importantly, two variables are created from the Sales Prices: A price-per-square-foot figure ("Last_Sale_Price") and a total Sale Price ("Last_Sale_Price_Total"). Sale Price per

Table 3.5: Sales By Asset Class

Bldg Code	Build Type	N	# Sales	Median Sale	Median Sale $/SF
A	One Family Dwellings	4435615	252283	$320,000	$215.85
B	Two Family Dwellings	3431762	219492	$340,000	$155.79
C	Walk Up Apartments	1873447	135203	$330,000	$67.20
D	Elevator Apartments	188689	45635	$398,000	$4.69
E	Warehouses	84605	5126	$200,000	$31.48
F	Factory	67174	4440	$350,000	$56.44
G	Garages	221620	13965	$0	$78.57
H	Hotels	10807	619	$5,189,884	$184.82
I	Hospitals	17650	687	$600,000	$62.66
J	Theatres	2662	152	$113,425	$4.01
K	Retail	265101	14841	$200,000	$60.63
L	Loft	18239	1259	$1,937,500	$101.36
M	Religious	78063	1320	$375,000	$91.78
N	Asylum	8498	190	$275,600	$35.90
O	Office	93973	5294	$550,000	$143.29
P	Public Assembly	15292	437	$350,000	$85.47
Q	Recreation	55193	232	$0	$0
R	Condo	78188	40157	$444,750	$12.65
S	Mixed Use Residence	467555	29396	$250,000	$78.29
T	Transportation	4012	49	$0	$0
U	Utility	32802	129	$0	$175
V	Vacant	449667	29091	$0	$134.70
W	Educational	38993	704	$0	$0
Y	Gov't	7216	44	$21,451.50	$0.30
Z	Misc	49583	2740	$0	$0

Table 3.6:   Sale Price Per Square Foot by Asset Class and Borough

Build Type	BK	BX	MN	QN	SI
Elevator Apartments	$2.65	$1.74	$10.80	$1.87	$1.23
Factory	$33.33	$53.19	$135.62	$92.42	$55.01
Garages	$78.94	$80.57	$94.43	$71.11	$67.46
Loft	$46.32	$78.26	$141.56	$150.37	$61.82
Office	$118.52	$123.04	$225.96	$148.45	$105
One Family Dwellings	$221.26	$176.98	$757.58	$232.69	$203.88
Two Family Dwellings	$140.95	$131.06	$296.10	$181.84	$160.76
Walk Up Apartments	$69.97	$84.05	$50.61	$36.94	$75.38

Table 3.7:   Number of Sales by Asset Class and Borough

Build Type	BK	BX	MN	QN	SI
Elevator Apartments	8,377	4,252	23,641	9,196	169
Factory	2,265	453	109	1,520	93
Garages	5,386	2,659	1,097	4,000	823
Loft	119	21	1,108	8	3
Office	1,112	340	2,081	1,162	599
One Family Dwellings	45,009	17,508	1,654	126,333	61,779
Two Family Dwellings	83,547	25,920	1,566	83,940	24,519
Walk Up Apartments	63,552	18,075	19,824	31,932	1,820

Table 3.8:   Base Modeling Data Features

Feature	Min	Median	Mean	Max
has_building_area	0	1.00	1.00	1.00
Percent_Com	0	0.00	0.16	1.00
Percent_Res	0	1.00	0.82	1.00
Percent_Office	0	0.00	0.07	1.00
Percent_Retail	0	0.00	0.04	1.00
Percent_Garage	0	0.00	0.01	1.00
Percent_Storage	0	0.00	0.02	1.00
Percent_Factory	0	0.00	0.00	1.00
Percent_Other	0	0.00	0.00	1.00
Last_Sale_Price	0	312.68	531.02	62,055.59
Last_Sale_Price_Total	2	2,966,835.00	12,844,252.00	1,932,900,000.00
Years_Since_Last_Sale	1	4.00	5.05	14.00
SMA_Price_2_year	0	296.92	500.89	62,055.59
SMA_Price_3_year	0	294.94	495.29	62,055.59
SMA_Price_5_year	0	300.12	498.82	62,055.59
Percent_Change_SMA_2	-1	0.00	685.69	15,749,999.50
Percent_Change_SMA_5	-1	0.00	337.77	6,299,999.80
EMA_Price_2_year	0	288.01	482.69	62,055.59
EMA_Price_3_year	0	283.23	471.98	62,055.59
EMA_Price_5_year	0	278.67	454.15	62,055.59
Percent_Change_EMA_2	-1	0.00	422.50	9,415,128.85
Percent_Change_EMA_5	-1	0.06	308.05	5,341,901.60

Square foot eventually becomes the outcome variable in one of the predictive models, even though it is referred to as Sale Price. Further features are derived which carry forward the previous sale price of a tax lot, if there is one, through successive years. Previous Sale Price is then used to create Simple Moving Averages (SMA), Exponential Moving Averages (EMA), and percent change measurements between the moving averages. In total, 69 variables are input to the feature engineering process and 92 variables are output. The final base modeling data set is 92 variables by 8,247,499 rows.

3.3.2   Zip Code Modeling Data

The first of the two comparative modeling data sets is the Zip Code modeling data. Using the base data as a starting point, several features are generated to describe characteristics of the Zip Code where each tax lot resides. A summary table of the Zip Code level features is presented in 3.9.

In general, the base model data features are aggregated to a Zip Code level and attached to the individual observations, including SMA and EMA calculations. Additionally, a second set of features are added, denoted as "bt_only", which filter only for tax lots of the same building type and aggregate to the Zip Code level. In total, the Zip Code feature engineering process inputs 92 variables and outputs 122 variables.

3.3.3   Spatial Lag Modeling Data

Spatial lags are variables created from physically proximate observations. For example, calculating the average age of all buildings within 100 meters of a tax lot constitutes a spatial lag. Creating spatial lags presents both advantages and disadvantages in the modeling process. Spatial lags allow for much more fine-tuned measurements of a building's surrounding area. Intuitively, knowing the average sale price of all buildings within 500 meters of a building can be more informative than knowing the sale prices of all buildings in the same Zip Code. However, creating spatial lags is computationally expensive. In addition, it can be difficult to

Table 3.9: Zip Code Modeling Data Features

Feature	Min	Median	Mean	Max
Last Year Zip Sold	0.00	27.00	31.14	112.00
Last Year Zip Sold Percent Ch	-1.00	0.00	Inf	Inf
Last Sale Price zip code average	0.00	440.95	522.87	1,961.21
Last Sale Price Total zip code average	10.00	5,312,874.67	11,877,688.55	1,246,450,000.00
Last Sale Date zip code average	12,066.00	13,338.21	13,484.39	17,149.00
Years Since Last Sale zip code average	1.00	4.84	4.26	11.00
SMA Price 2 year zip code average	34.31	429.26	501.15	2,092.41
SMA Price 3 year zip code average	34.31	422.04	496.47	2,090.36
SMA Price 5 year zip code average	39.48	467.04	520.86	2,090.36
Percent Change SMA 2 zip code average	-0.20	0.04	616.47	169,999.90
Percent Change SMA 5 zip code average	-0.09	0.03	341.68	113,333.27
EMA Price 2 year zip code average	30.77	401.43	479.38	1,883.81
EMA Price 3 year zip code average	33.48	419.11	479.95	1,781.38
EMA Price 5 year zip code average	29.85	431.89	472.80	1,506.46
Percent Change EMA 2 zip code average	-0.16	0.06	388.90	107,368.37
Percent Change EMA 5 zip code average	-0.08	0.07	326.17	107,368.38
Last Sale Price bt only	0.00	357.71	485.97	6,401.01
Last Sale Price Total bt only	10.00	3,797,461.46	11,745,130.56	1,246,450,000.00
Last Sale Date bt only	12,055.00	13,331.92	13,497.75	17,149.00
Years Since Last Sale bt only	1.00	4.78	4.30	14.00
SMA Price 2 year bt only	0.00	347.59	462.67	5,519.39
SMA Price 3 year bt only	0.00	345.40	458.50	5,104.51
SMA Price 5 year bt only	0.00	372.30	481.09	4,933.05
Percent Change SMA 2 bt only	-0.55	0.03	600.10	425,675.69
Percent Change SMA 5 bt only	-0.33	0.02	338.15	188,888.78
EMA Price 2 year bt only	0.00	332.98	442.79	5,103.51
EMA Price 3 year bt only	0.00	332.79	443.02	4,754.95
EMA Price 5 year bt only	0.00	340.57	436.70	4,270.37
Percent Change EMA 2 bt only	-0.47	0.06	377.17	254,462.97
Percent Change EMA 5 bt only	-0.34	0.06	335.17	178,947.30

set a proper radius for the spatial lag calculation; in a city, 500 meters may be appropriate (for certain building types), whereas several kilometers or more may be appropriate for less densley populated areas. In this paper, we present a solution for the computational challenges and suggest a potential approach to solving the radius-choice problem.

3.3.3.1 Creating the Point-Neighbor Relational Graph

To build our spatial lags, for each point in the data, we must identify which of all other points in the data fall within a specified radius. This requires iteratively running point-in-polygon operations, i.e., "given polygon P and an arbitrary point q, determine whether point q is enclosed by the edges of the polygon" (Huang 1996). This process is conceptually illustrated in figure 3.2.
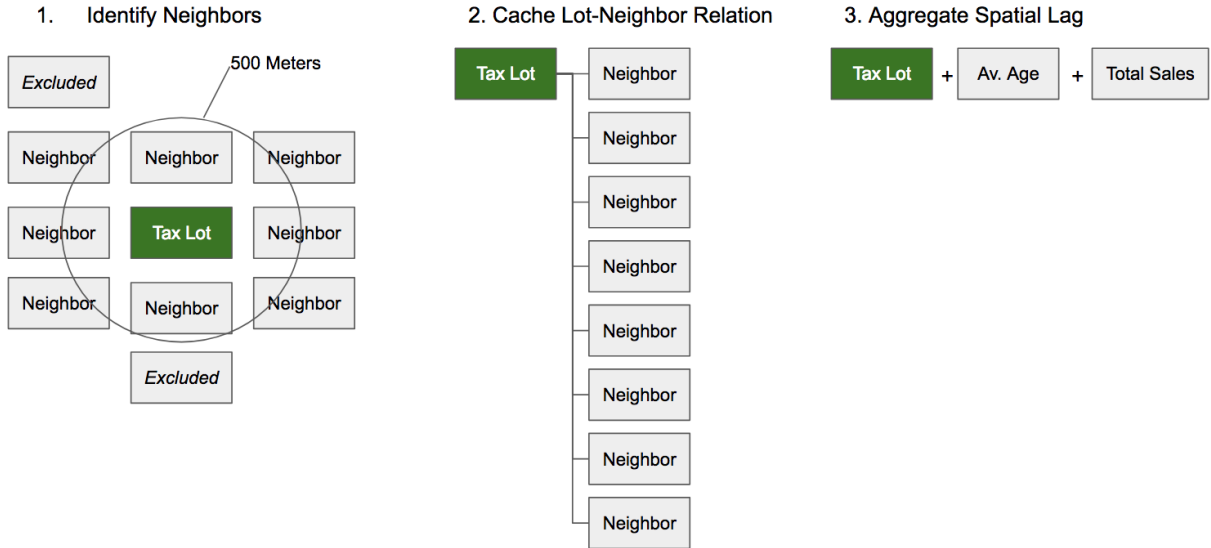

Figure 3.2: Spatial Lag Feature Creation Process

Given that, for every point q_i in our dataset, we need to determine whether every other point q_i falls within a given radius, this means that the time-complexity of our operation can be approximated as:

$$O(N(N-1))$$

Since the number of operations approaches N^2, calculating spatial lags for all 8,247,499 observations in our modeling data would be infeasible from a time and computation perspective. Assuming that tax lots rarely if ever move over time, we first reduced the task to the number of unique tax lots in New York City from 2003-2017, which is 514,124 points. Next, we implemented an indexing technique that greatly speeds up the process of creating a point-neighbor relational graph. The indexing technique both reduces the relative search space for each computation and also allows for parallelization of the point-in-polygon operations by dividing the data into a gridded space. The gridded spatial indexing process is outlined in Algorithm 1.

Algorithm 1 Gridded Spatial Indexing

1: **for** each grid partition G **do**
2:    Extract all points points G_i contained within partition G
3:    Calculate convex hull $H(G)$ such that the buffer extends to distance d
4:    Define Search space S as all points within Convex hull $H(G)$
5:    Extract all points S_i contained within S
6:    **for** each data point G_i **do**
7:       Identify all points points in S_i that fall within $abs(G_i + d)$
8:    **end for**
9: **end for**

Each gridded partition of the data is married with a corresponding search space S, which is the convex hull of the partition space buffered by the maximum distance d. In our case, we are buffering the search space by 500 meters. Choosing an appropriate radius for buffering presents an additional challenge in creating spatially-concious machine learning predictive models. In this paper, we choose an arbitrary radius, and use a two-stage modeling process to test the appropriateness of that assumption. In future work, implementing an "adaptive bandwidth" technique using cross-validation to determine the optimal radius could be done.

By partitioning the data into spatial grids, we are able to reduce the search-space for each operation by an arbitrary number of partitions G. This improves the base run-time complexity to:

$$O(N(\frac{N-1}{G})$$

By making G arbitrarily large (bounded by computational resources only), we reduced runtime substantially. Furthermore, binning the operations into grids allowed us to parallelize the computation, further reducing the overall run time. Figure 3.3 shows a comparison of computation times between the basic point-in-polygon technique and a sequential version of the grided indexing technique. Note that the grid method starts out as slower than the basic point-in-polygon technique due to pre-processing overhead, but wins out in terms of speed as complexity increases. This graph also does not reflect parallelization of the operation, which further reduces the time required to calculate the point-neighbor relational graph.


Figure 3.3: Spatial Index Time Comparison

3.3.3.2 Calculating Spatial Lags

Once the origin-neighbor relational graph has been constructed, we can aggregate the data into spatial lag variables. One advantage of using spatial lags is the rich number of potential features which can be engineered. Spatial lags can be weighted based on a distance

function, e.g., physically closer observations can be given more weight. For our modeling purposes, we created two sets of features: distance weighted features (denoted with a "_dist" in Table 3.10) and simple average features (denoted with "_basic" in Table 3.10). SMA and EMA as well as percent changes were also calculated.

Temporal and spatial derivatives of the Spatial Lag features presented in Table 3.10 were also added to the model, including: variables weighted by euclidean distance ("dist"), basic averages of the spatial lag radius ("basic mean"), Simple Moving Averages ("SMA") for 2 years, 3 years and 5 years, exponential moving averages ("EMA") for 2 years, 3 years and 5 years, and year-over-year percent changes for all variables ("perc change"). In total, the spatial lag feature engineering process input 92 variables and output 194 variables. A summary of the Spatial Lag features are presented in Table 3.10.

3.4 Outcome Variables

The final step in creating the modeling data is to define the outcome variables. The outcome variables reflect the prediction tasks; a binary variable for classification and a continuous variable for regression:

1) **Binary: Sold** whether a tax lot sold in a given year. Used in the Probability of Sale classification model.

2) **Continuous: Sale Price per SF** The price-per-square foot associated with a transaction, if a sale took place. Used in the Sale Price Regression model.

Table 3.11 describes the distributions of both outcome variables.

3.5 Algorithms Comparison

We implemented and compared several algorithms across our two-stage process. In Stage 1, the Random Forrest algorithm was used to identify the optimal subset of building types and geographies for our spatial lag aggregation assumptions. In Stage 2, we analyze the performance of the Random Forrest algorithm again, as well as Generalized Linear Model

Table 3.10:   All Spatial Lag Features

Feature	Min	Median	Mean	Max
Radius_Total_Sold_In_Year	1.00	20.00	24.00	201.00
Radius_Average_Years_Since_Last_Sale	1.00	4.43	4.27	14.00
Radius_Res_Units_Sold_In_Year	0.00	226.00	289.10	2,920.00
Radius_All_Units_Sold_In_Year	0.00	255.00	325.94	2,923.00
Radius_SF_Sold_In_Year	0.00	259,403.00	430,891.57	8,603,639.00
Radius_Total_Sold_In_Year_sum_over_2_years	2.00	41.00	48.15	256.00
Radius_Average_Years_Since_Last_Sale_sum_over_2_years	2.00	9.25	8.70	26.00
Radius_Res_Units_Sold_In_Year_sum_over_2_years	0.00	493.00	584.67	3,397.00
Radius_All_Units_Sold_In_Year_sum_over_2_years	1.00	555.00	660.67	4,265.00
Radius_SF_Sold_In_Year_sum_over_2_years	2,917.00	580,947.00	872,816.44	14,036,469.00
Radius_Total_Sold_In_Year_percent_change	-0.99	0.00	0.27	77.00
Radius_Average_Years_Since_Last_Sale_percent_change	-0.91	0.13	0.26	8.00
Radius_Res_Units_Sold_In_Year_percent_change	-1.00	-0.04	Inf	Inf
Radius_All_Units_Sold_In_Year_percent_change	-1.00	-0.04	Inf	Inf
Radius_SF_Sold_In_Year_percent_change	-1.00	-0.02	Inf	Inf
Radius_Total_Sold_In_Year_sum_over_2_years_percent_change	-0.96	-0.03	0.03	15.00
Radius_Average_Years_Since_Last_Sale_sum_over_2_years_percent_change	-0.72	0.12	0.17	2.50
Radius_Res_Units_Sold_In_Year_sum_over_2_years_percent_change	-1.00	-0.04	Inf	Inf
Radius_All_Units_Sold_In_Year_sum_over_2_years_percent_change	-0.99	-0.04	0.12	84.00
Radius_SF_Sold_In_Year_sum_over_2_years_percent_change	-0.98	-0.04	0.18	361.55
Percent_Com_dist	0.00	0.04	0.07	0.56
Percent_Res_dist	0.00	0.46	0.43	0.66
Percent_Office_dist	0.00	0.01	0.03	0.48
Percent_Retail_dist	0.00	0.02	0.02	0.09
Percent_Garage_dist	0.00	0.00	0.00	0.27
Percent_Storage_dist	0.00	0.00	0.01	0.26
Percent_Factory_dist	0.00	0.00	0.00	0.04
Percent_Other_dist	0.00	0.00	0.00	0.09
Percent_Com_basic_mean	0.00	0.04	0.07	0.54
Percent_Res_basic_mean	0.00	0.46	0.43	0.66
Percent_Office_basic_mean	0.00	0.01	0.03	0.44
Percent_Retail_basic_mean	0.00	0.02	0.02	0.08
Percent_Garage_basic_mean	0.00	0.00	0.00	0.29
Percent_Storage_basic_mean	0.00	0.00	0.01	0.23
Percent_Factory_basic_mean	0.00	0.00	0.00	0.03
Percent_Other_basic_mean	0.00	0.00	0.00	0.04
Percent_Com_dist_perc_change	-0.90	0.00	0.00	6.18
Percent_Res_dist_perc_change	-0.50	0.00	0.03	36.73
Percent_Office_dist_perc_change	-1.00	0.00	Inf	Inf
Percent_Retail_dist_perc_change	-0.82	0.00	Inf	Inf
Percent_Garage_dist_perc_change	-1.00	0.00	Inf	Inf
Percent_Storage_dist_perc_change	-1.00	-0.01	Inf	Inf
Percent_Factory_dist_perc_change	-1.00	0.00	Inf	Inf
Percent_Other_dist_perc_change	-1.00	0.00	Inf	Inf
SMA_Price_2_year_dist	0.00	400.01	496.30	3,816.57
SMA_Price_3_year_dist	0.00	396.94	492.00	3,816.57
SMA_Price_5_year_dist	8.83	425.55	515.29	3,877.53
Percent_Change_SMA_2_dist	-0.13	0.03	552.33	804,350.67
Percent_Change_SMA_5_dist	-0.09	0.02	317.46	322,504.58
EMA_Price_2_year_dist	0.00	378.63	475.54	3,431.17
EMA_Price_3_year_dist	8.83	382.25	476.05	3,296.46
EMA_Price_5_year_dist	7.88	386.34	468.91	2,813.34
Percent_Change_EMA_2_dist	-0.09	0.06	346.51	480,829.57
Percent_Change_EMA_5_dist	-0.02	0.06	303.55	273,458.42
SMA_Price_2_year_basic_mean	0.02	412.46	496.75	2,509.79
SMA_Price_3_year_basic_mean	0.02	409.00	492.43	2,509.79
SMA_Price_5_year_basic_mean	17.16	443.34	515.67	2,621.01
Percent_Change_SMA_2_basic_mean	-0.13	0.04	543.51	393,749.99
Percent_Change_SMA_5_basic_mean	-0.09	0.03	312.46	157,500.00
EMA_Price_2_year_basic_mean	0.02	390.30	475.96	2,259.21
EMA_Price_3_year_basic_mean	11.39	393.25	476.45	2,136.36
EMA_Price_5_year_basic_mean	15.30	402.06	469.09	1,848.27
Percent_Change_EMA_2_basic_mean	-0.09	0.06	340.89	235,378.24
Percent_Change_EMA_5_basic_mean	-0.02	0.06	296.78	133,547.59

Table 3.11: Distributions for Outcome Variables

	Sold	Sale Price per SF
Min.	0.00	0.0
1st Qu.	0.00	163.5
Median	0.00	375.2
Mean	0.04	644.8
3rd Qu.	0.00	783.3
Max.	1.00	83,598.7

(GLM), Gradient Boosting Machines (GBM), and a Feed-Forward Artificial Neural Network (ANN). Each algorithm is appropriate for both classification and regression tasks.

3.5.1 Random Forrest

The Random Forest concept was proposed by Leo Breiman in 2001 as the ensembling of prediction decision trees iteratively trained across randomly generated subsets of data (Leo Breiman 2001). Algorithm 2 outlines the procedure (Hastie, Tibshirani, and Friedman 2001).

Algorithm 2 Random Forest for Regression or Classification

1. For $b = 1$ to B
   - (a) Draw a bootstrap sample Z of the size N from the training data.
   - (b) Grow a random-forest tree T_b to the bootstrapped data, by recursively repeating the following steps for each terminal node of the tree, until the minimum node size n_{min} is reached.
     - i. Select m variables at random from the p variables
     - ii. Pick the best variable/split-point among the m.
     - iii. Split the node into two daughter nodes.
2. Output the ensemble of trees $\{T_b\}_1^B$.

To make a prediction at a new point x:

Regression: $\hat{f}_{rf}^B(x) = \frac{1}{B} \sum_{b=1}^{B} T_b(x)$

Classification: Let $\hat{C}_b(x)$ be the class prediction of the bth random-forest tree. Then $\hat{C}_{rf}^B(x) = $ *majority vote* $\{\hat{C}_b(x)\}_1^B$

Previous works (see: Antipov and Pokryshevskaya (2012); also Schernthanner H. (2016)) have found the Random Forest algorithm (L. Breiman 2001) suitable to prediction tasks involving real estate. While algorithms exist that may marginally outperform Random Forest

in terms of predictive accuracy (such as neural networks and functional gradient descent algorithms), Random Forest is highly scalable and parallelizable, and therefore a natural choice for quickly assessing the predictive power of different feature engineering techniques. For these reasons and more outlined below, we selected Random Forrest as the primary algorithm for Stage 1 of our modeling process.

Random Forest can be used for both classification and regression tasks. The Random Forest algorithm works by generating a large number of independent classification or regression decision trees and then employing majority voting (for classification) or averaging (for regression) to generate predictions. Over a data set of N rows by M predictors, a bootstrap sample of the data is chosen (n < N) as well as a subset of the predictors (m < M). Individual decision/regression trees are built on the n by m sample. Because the trees can be built independently (and not sequentially, as is the case with most functional gradient descent algorithms), the tree building process can be executed in parallel. With a sufficiently large number of cores, the model training time can be significantly reduced.

We chose Random Forrest as the algorithm for Stage 1 because:

1) The algorithm can be parallelized and is relatively fast compared to neural networks and functional gradient descent algorithms

2) Can accommodate categorical variables with many levels. Real estate data often contains information describing the location of the property, or the property itself, as one of a large set of possible choices, such as neighborhood, county, census tract, district, property type, and zoning information. Because factors need to be recoded as individual dummy variables in the model building process, factors with many levels will quickly encounter the curse of dimensionality in multiple regression techniques.

3) Appropriately handles missing data. Predictions can be made with the parts of the tree which are successfully built, and therefore, there is no need to filter out incomplete observations or impute missing values. Since much real estate data is self reported, incomplete fields are common in the data.

4) Robust against outliers. Because of bootstrap sampling, outliers appear in individual trees less often, and therefore, are reduced in terms of importance. Real estate data, especially with regards to pricing, tends to contain outliers. For example, the dependent variable in one of our models, Sale Price, shows a clear divergence in median and mean, as well as a maximum significantly higher than the third quartile.

5) Can recognize non-linear relationships in data, which is useful when modeling spatial relationships.

6) Is not affected by co-linearity in the data. This is highly valuable as real estate data can be highly correlated.

To run the model, we have chosen the h2o.randomForest implementation from the h2o R open source library. The h2o implementation of the Random Forest algorithm is particularly well-suited for high parallelization. For more information, see: https://www.h2o.ai/.

3.5.2   Generalized Linear Model (GLM)

3.5.3   Gradient Boosting Machine

3.5.4   Feed-Forward Artificial Neural Netwrok (ANN)

3.6   Model Validation

The goal of the predictive models are to be able to successfully predict both the probability and amount of real estate sales into the near future. As such, our models will use out-of-time validation to assess performance. As shown in Figure 3.4 The models will be trained using data from 2003-2015. 2016 modeling data will be used during the model training process as cross-validation data. Finally, we will score our model using 2017 as a hold-out sample. Using out-of-time validation should ensure that the models generalize well into the immediate future.

Figure 3.4: Out-of-time validation

3.7 Variable Selection

For ease of processing and to improve the ability of the model to generalize into the future, a variable selection step is added to the modeling process. A Random Forest model is first trained on a 1% sub-sample of the modeling data. Variable importance of the resulting model is calculated using the technique proposed by Friedman (2001), i.e., for a collection of decision trees $[T_m]_1^m$:

$$\hat{I}_j^2 = \frac{1}{M} \sum_{m=1}^{M} \hat{I}_j^2(T_m)$$

Where influence I for variable j is calculated as the sum of corresponding improvements in squared-error for node tree T. After calculating variable importance for the model data subset, the variables are rank-ordered by descending importance. Variables which account for 80% of the total variable importance are chosen to advance to the model training round on the full modeling data sets.

3.8 Evaluation Metrics

We have chosen evaluation metrics that will allow us to easily compare the performance of the models against other models with the same outcome variable. The classification models (Probability of Sale) will be compared using Area Under the ROC Curve (AUC). The

regression models (Sale Price) will be compared using Root Mean Squared Error (RMSE). Both evaluation metrics are common for their respective outcome variable types, and as such will be useful for comparing within model-groups.

3.8.1   Area Under ROC Curve (AUC)

A classification model typically outputs a probability that a given row in the data belongs to a group. In the case of binary classification, the value falls between 0 and 1. There are many techniques for determining the cut off threshold for classification; a typical method is to assign anything above a 0.5 into the "1" or positive class. An ROC curve (receiver operating characteristic curve) plots the True Positive Rate vs. the False Positive rate at different classification thresholds; it is a measurement of the performance of a classification model across all possible thresholds, and therefore sidesteps the need to arbitrarily assign a cutoff.

Area Under the ROC Curve, or AUC measures the entire two-dimensional area underneath the ROC curve. It is the integration of the curve from (0,0) to (1,1), defined as $AUC = \int_{(0,0)}^{(1,1)} f(x)dx$.

AUC provides a relatively standard measure of performance across all possible classification thresholds, and can be interpreted as the probability that the model ranks a random positive example more highly than a random negative example. A value of 0.5 represents a perfectly random model, while a value of 1.0 represents a model that can perfectly discriminate between the two classes. AUC is useful for comparing classification models against one another because they are both scale and threshold-invariant.

One of the drawbacks to AUC is that is does not describe the trade-offs between false positives and false negatives. In certain circumstances, a false positive might be considerably less desirable than a false negative, or vice-versa. For our purposes, we rank false positives and false negatives as equally undesirable outcomes.

3.8.2   Root Mean Squared Error

The Root Mean Squared Error (RMSE) is a common measurement of the differences between values predicted by a regression model and the observed values. It is formally defined as $RMSE = \sqrt{\frac{\sum_1^T (\hat{y}_t - y_t)^2}{T}}$, where $\hat{y}$ represents the prediction and y represents the observed value at observation t.

Lower RMSE scores are typically more desirable. An RMSE value of 0 would indicate a perfect fit to the data. RMSE can be difficult to interpret on its own, however, it is useful for comparing models with similar outcome variables. In our case, the outcome variables (Sales Price per Square Foot) are consistent across modeling data sets, and therefore can be reasonably compared using RMSE.

4   Results

4.1   Summary of Results

We have conducted comparative analyses across a two-stage modeling proces. In Stage 1, using the Random Forrest algorithm, we tested 3 competing feature engineering techniques (base, zip-code aggregation and spatial-lag aggreation) for both a classification task (predicting the occurance of a building sale) and a regression task (predicting the sale price of a building). We analyzed the results of the first stage to identify which geographies and building types our model assumptions worked best. In Stage 2, using a subset of the modeling data (selected via an analysis of the output from Stage 1), we compared four algorithms–Generalized Linear Model (GLM), Random Forrest (RF), Gradient Boosted Machine (GBM) and Feed-Forward Multilayer Artificial Neural Network (ANN)–across our 3 competing feature engineering techniques for both classification and regression tasks. We analyzed the performance of the different model/data combos as well as conducted an analysis of the variable importances for the top performing models.

In Stage 1 (Random Forrest, using all data), we found that models which utilized spatial features outperformed those models using zip-code features the majority of the time for both classification and regression. Of three models, the Sale Price Regression model using Spatial features finished 1st or 2nd 24.1% of the time (using Root Mean Squared Error as a ranking criterion), while the Zip Code Regression model finished in the top two spots only 11.2% of the time. Both models performed worse than the Base Regression model overall, which ranked in 1st or 2nd place 31.5% of the time. The story for the classification models was largely the same: the Spatial features tended to outperform the Zip Code data while the Base data won out overall. All models had similar performances on training data, but the Spatial and Zip Code datasets tended to underperform when generalizing to the hold-out test data, suggesting problems with overfitting.

We then analyzed the performance of both the regression and classification Random Forrest models by geography and building type. We found that the models performed considerably better on Walk Up Apartments and Elevator Buildings (building types C and D) and in Manhattan, Brooklyn and the Bronx. Using these as filtering criteria, we created a subset of the modeling data for the subsequent modeling stage.

During Stage 2 (many algorithms using a subset of modeling data), we compared four algorithms across the same three competing feature engineering techniques using a filtered subset of the original modeling data. Unequivocally, the spatial features performed best across all models and tasks. For the classification task, the GBM algorithms performed best in terms of AUC, followed by ANN and Random Forrest. For regression, the ANN algorithms performed best (as measured by Mean Absolute Error, Root Mean Squared Error and R-squared) with the spatial features ANN model performing best.

We conclude that spatial lag features can significantly increase the accuracy of machine learning-based real estate sale prediction models. We find that model overfitting presents a challenge when using spatial features, but that this can be overcome by implementing different algorithms, specifically ANN and GBM. Finally, we find that our implementation of

Table 4.1:   Sale Price Model Rankings, RMSE by Borough and Building Type

Model Rank	1	2	3	Average Rank
Base	22.2%	9.3%	1.9%	1.39
Spatial Lag	5.6%	18.5%	9.3%	2.11
Zip	5.6%	5.6%	22.2%	2.50

Table 4.2:   Sale Price Model RMSE For Validation and Test Hold-out Data

type	base	zip	spatial lag
Validation	280.63	297.97	286.23
Test	287.83	300.60	297.92

spatial-lag features works best for certain kinds of buildings in specific geographic areas, and we hypothesize that this is due to the assumptions made when building the spatial features.

4.2   Stage 1) Random Forrest Models Using All Data

4.2.1   Results of Random Forrest Sale Pice Regression Model

We analyzed the the Root Mean Squared Errors (RMSE) of the Random Forrest models predicting Sale Price across feature engineering methods, Borough and Building Type. Table 4.1 displays the average ranking by model type as well as the distribution of models that ranked first, second and third for each respective Borough/Building Type combination. When we rank the models by performance for each Borough, Building Type combination, we find that the Spatial Lag models outperform the Zip Code models in 72% of cases with an average model-rank of 2.11 and 2.5, respectively.

The Base modeling data set tends to outperform both enriched datasets, suggesting an issue with model overfitting in some areas. We see further evidence of overfit in Table 4.2 where, despite similar performances on Validation data, the Zip and Spatial models have higher validation-to-test-set spreads. Despite this, the Spatial Lag features outperform all other models in certain locations, notably in Manhattan as shown in Figure 4.1.

Figure 4.1: RMSE By Borough and Building Type

Figure 4.1 displays test RMSE by model, faceted by Borough on the y-axis and Building Type on the x-axis (See Table 3.3 and Table 3.5 for a description of building type codes). We make the following observations from Figure 4.1:

- The Spatial modeling data outperforms both Base and Zip Code in 6 cases, notably for Type A buildings (One Family Dwellings) and Type L buildings (Lofts) in Manhattan as well as Type O Buildings (Office) in Queens

Table 4.3: Probability of Sale Model AUC

Model AUC	Base	Zip	Spatial Lag
Validation	0.832	0.829	0.829
Test	0.830	0.825	0.828

- The "residential" building Types A (One Family Dwellings), B (Two Family Dwellings), C (Walk Up Apartments) and D (Elevator Apartments) have generally lower RMSE scores compared to the non-residential types

- Spatial features perform best in Brooklyn, Bronx and Manhattan and for residential building types

4.2.2 Results of Random Forrest Probability of Sale Classification Model

Similar to the results of the Sale Price regression models we find the Spatial models perform better on the hold-out test data compared to the Zip Code data when using Area Under the ROC Curve (AUC) as an evaluation metric, as shown in Table 4.3. The Base Modeling data continues to outperform the Spatial and Zip Code data overall.

Figure 4.2 shows a breakdown of model AUC faceted along the x-axis by Building Type and along the y-axis by Borough. The coloring indicates by how much a model's AUC diverges from the cell average, which is useful for spotting overperformers. We make the following observations about Figure 4.2:

- The Spatial models outperform all other models for Elevator Buildings (Type D) and Walk Up Apartments (Type C), particularly in Brooklyn, the Bronx and Manhattan

- Classification tends to perform poorly in Manhattan vs. other Boroughs

- The Spatial models performs well in Manhattan for the residential building types (A, B, C and D)

Figure 4.2: AUC By Borough and Building Type

If we rank the classification models' performance for each Borough and Building Type, we see that the Spatial models consistently outperform the Zip Code models, as shown in Table 4.4. From this (as well as from similar patterns seen in the regression models) we can infer that the Spatial data is a superior data engineering technique, however, the algorithm used needs to account for potential model overfitting. In the next section, we discuss refining

Table 4.4:   Distribution and Average Model Rank for Probability of Sale by AUC across Borough and Building Types

Model Rank	1	2	3	Average Rank
Base	16.2%	12.0%	5.1%	2.22
Spatial Lag	11.1%	13.7%	8.5%	2.09
Zip	6.0%	7.7%	19.7%	1.69

the data used as well as employing different algorithms to maximize predictive capability of the Spatial features.

4.3   Model Comparisons Using Specific Geographies and Building Types

Using the results from the first modeling exercise, we conclude that Walk Up Apartments and Elevator Buildings in Manhattan, Brooklyn and the Bronx are suitable candidates for prediction using our current assumptions. These buildings share the characteristics of being residential as well as being fairly uniform in their geographic density. We analyze the performance of four algorithms (GLM, RF, GBM and ANN), using three feature engineering techniques, for both classification and regression, making the total number `4 x 3 x 2 = 24` models.

4.3.1   Regression Model Comparisons

The following criteria, described in detail in the Methodology section, were used to assess the predictive accuracy of the various regression models: Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), Mean Squared Error (MSE) and R-Squared. These four indicators are calculated using the hold-out test data, which allows us to ensure that the models perform well when predicting sale prices into the near future. The comparison metrics are presented in Table 4.5 and Figure 4.3.

The following conclusions can be made from Table 4.5 and Figure 4.3:

Table 4.5: Prediction Accuracy of Regression Models on Test Data

Data	Model	RMSE	MAE	MSE	R2
1) Base	GLM	446.35	221.16	199227.6	0.12
2) Zip	GLM	426.93	206.49	182270.1	0.19
3) Spatial	GLM	382.32	195.00	146170.5	0.35
1) Base	RF	387.99	174.24	150536.3	0.33
2) Zip	RF	475.20	190.33	225811.7	0.00
3) Spatial	RF	430.92	180.17	185695.5	0.18
1) Base	GBM	384.11	179.27	147543.5	0.35
2) Zip	GBM	454.53	186.00	206593.1	0.09
3) Spatial	GBM	406.70	170.97	165408.0	0.27
1) Base	ANN	363.02	178.58	131782.5	0.42
2) Zip	ANN	360.88	171.22	130232.2	0.42
3) Spatial	ANN	337.94	158.91	114202.0	0.49

1) The feed-forward Multilayer Artificial Neural Networks (ANN) perform best in nearly every metric across nearly all feature sets, with Gradient Boosted Machine (GBM) a close second in some circumstances

2) ANN and Generalized Linear Models (GLM) improve linearly in all metrics as you move from Base to Zip to Spatial, with Spatial performing the best. GBM and Random Forrest, on the other hand, perform best on the Base and Spatial feature sets and poorly on the Zip features

3) The highest model R-Squared is the ANN using Spatial features at 0.494, indicating that this model can account for nearly 50% of the variance in the test data

4) We see a similar pattern in the Random Forrest results compared to the previous modeling exercise using the full dataset: Base Features outperforming both Spatial and Zip, with Spatial coming in second consistently. This further validates our reasoning that Spatial features are highly predictive but suffer from overfitting and other algorithm-related reasons
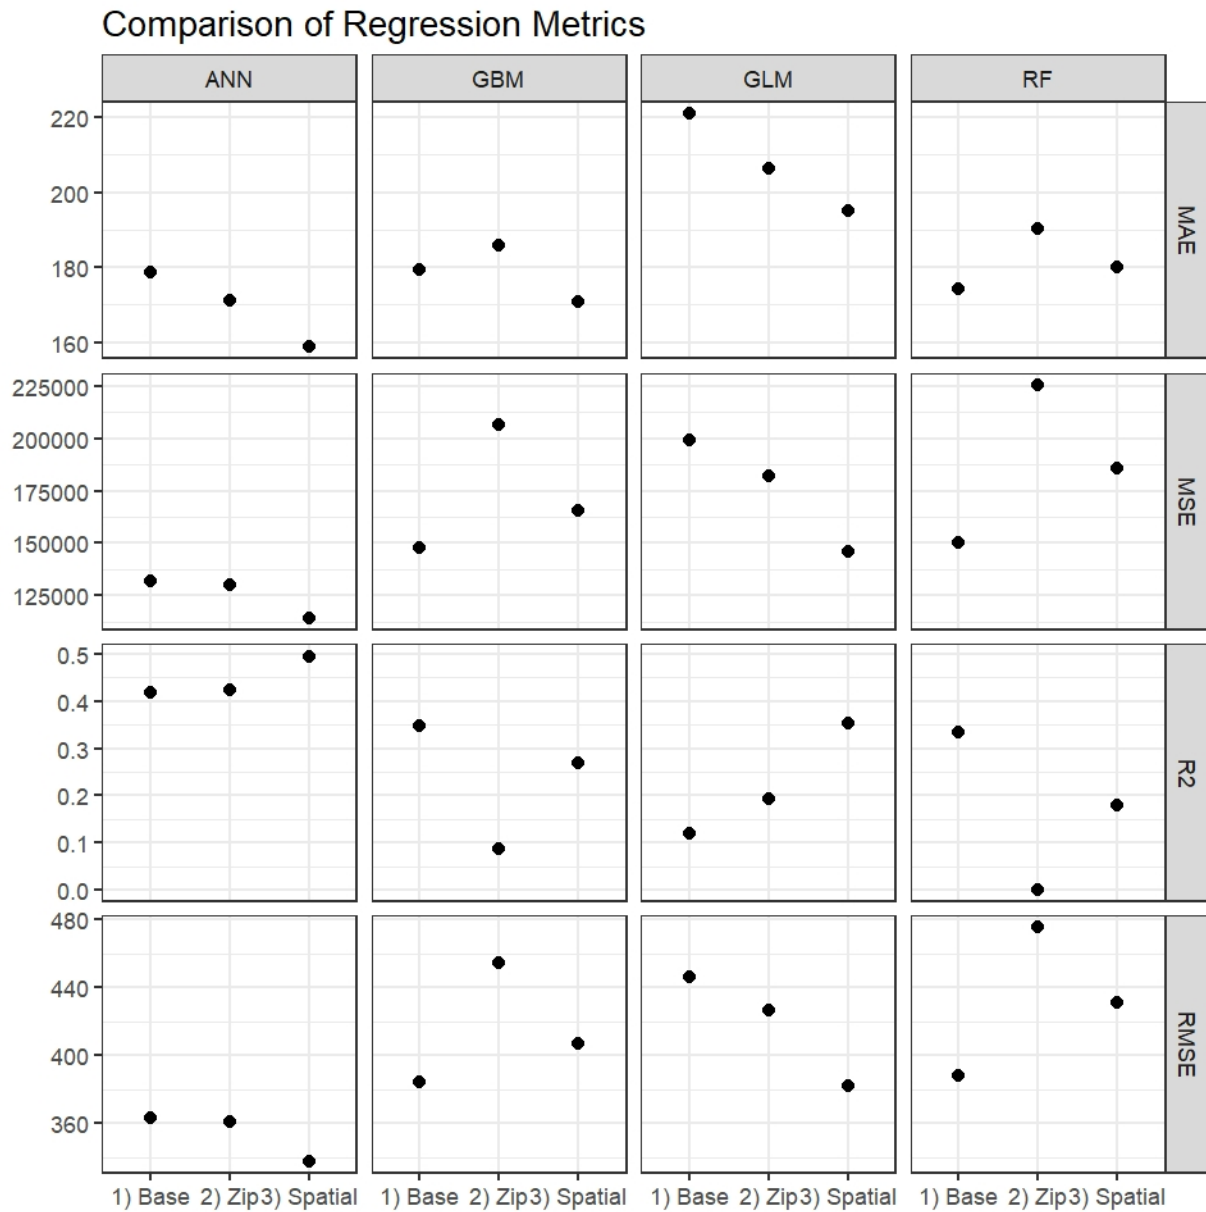
Figure 4.3: Comparative Regression Metrics

Figure 4.4 shows clusters of performance across R-Squared and Mean Absolute Error, with the ANN models outperforming their peers. This figure also makes clear that the marriage of Spatial features with the ANN algorithm results in a dramatic reduction in error rate compared to the other techniques.

Figure 4.4: Regression Model Performances On Test Data

4.3.2 Classification Model Comparisons

The criteria used to assess the accuracy of the classification models are as follows: Area Under the ROC Curve (AUC), Mean Squared Error (MSE), Root Mean Squared Error (RMSE) and R-squared (R2). As with the regression models, these four indicators are calculated using the hold-out test data, allowing us to ensure that the models generalize well

Table 4.6:   Prediction Accuracy of Classification Models on Test Data

Data	Model	AUC	MSE	RMSE	R2
1) Base	GLM	0.57	0.03	0.17	0.00
2) Zip	GLM	0.58	0.03	0.17	0.00
3) Spatial	GLM	0.50	0.03	0.17	-0.01
1) Base	RF	0.58	0.03	0.17	-0.03
2) Zip	RF	0.56	0.03	0.17	-0.06
3) Spatial	RF	0.78	0.03	0.17	0.00
1) Base	GBM	0.61	0.03	0.17	-0.03
2) Zip	GBM	0.61	0.03	0.17	-0.03
3) Spatial	GBM	0.82	0.03	0.16	0.04
1) Base	ANN	0.55	0.03	0.17	-0.03
2) Zip	ANN	0.57	0.03	0.17	-0.04
3) Spatial	ANN	0.76	0.03	0.17	-0.01

into the near future. The comparison metrics are presented in Table 4.6. Figure 4.5 shows the ROC curves and corrsponding AUC for each algorithm/feature set combination.

The following conclusions can be made from Table 4.6 and Figure 4.5:

1) Unlike the regression models, the GBM algorithm with Spatial features proved to be the best performing classifier. All Spatial models performed relatively well with the exception of the GLM Spatial model

2) The error metrics tend to be quite low and frequently negative given that they are calculated using probabilitites and classes (0,1)

3) Only 3 models have positive R-Squared values: ANN Spatial, RF Spatial and GLM Base. This would indicate that these models are adept at prediciting positive cases (occurances of sales) in the test data

4) GLM Spatial returned an AUC of less than 0.5, indicating a model that is conceptually worse than random. This is likely a result of extreme overfitting

Figure 4.5: Comparison of Classification Model ROC Curves

Figure 4.6 plots the individual models by AUC and R-Squared. The Spatial models tend to outperform the other models by a significant margin. Interestingly, when compared to the regression model scatterplot, Figure 4.4, the classification models tend to cluster in their performance by feature set. In 4.4, we see the regression models tending to cluster by algorithm rather than features.

Figure 4.6: Scatterplot of Classification Models

4.4   Variable Importance Analysis of Top Performing Models

Feature importance is calculated for each algorthm as being proportional to the average decrease in the squared error after including that variable in the model. The most important variable gets a score of 1; scores for other variables are derived by standardizing their measured reduction in error relative to the largest one. The top 10 variables for both the most successful

Table 4.7:   Feature Importance of Top Performing Regression Model

Variable	Description	Scaled Importance (Max = 1)	Cumulative %
BuiltFAR	Floor area ratio built	1.000	1.80%
FacilFAR	Maximum Allowable Floor Area Ratio	0.922	3.40%
Last_Sale_Price_Total	The previous sale price	0.901	5.10%
Last_Sale_Date	Date of last sale	0.893	6.70%
Last_Sale_Price	The previous sale price	0.870	8.20%
Years_Since_Last_Sale	Number of years since last sale	0.823	9.70%
ResidFAR	Floor Area Ratio not yet built	0.814	11.20%
lon	Longitude	0.773	12.60%
Year	Year of record	0.759	13.90%
BldgDepth	Square feet from font to back	0.758	15.30%

Table 4.8:   Feature Importance of Top Performing Classification Model

Variable	Description	Scaled Importance (Max = 1)	Cumulative %
Percent_Neighbords_Sold	Percent of Nearby Properties Sold in the Previous Year	1.000	21.90%
Percent_Office	Percent of the build which is Office	0.698	37.20%
Percent_Garage	Percent of the build which is Garage	0.634	51.10%
Percent_Storage	Percent of the build which is Storage	0.518	62.40%
Building_Age	The Age of the building	0.225	67.40%
Last_Sale_Price	Price of building last time is was sold	0.165	71.00%
Percent_Retail	Percent of the build which is Retail	0.147	74.20%
Years_Since_Last_Sale	Year since building last sold	0.121	76.90%
ExemptTot	Total tax exempted value of the building	0.069	78.40%
Radius_Res_Units_Sold_In_Year	Residential units within 500 meters sold in past year	0.056	79.60%

regression and most successful classification models are presented in tables 4.7 and 4.8.

We observe that the regression model has a much higher dispersion of feature importances compared to the classification model. The top variable in the regression model, BuiltFAR, which is a measure of how much of a building's floor to area ratio has been used (a proxy for overall building size) contributing only 1.8% of the reduction in error rate in the overall model. Conversely, in the classification model, we see the top variable, "Percent_Neighbors_Sold" (a measure of how many buildings within 500 meters were sold in the past year) contributes 21.9% of the total reduction in squared error.

Variable importance analysis of the regression model indicates that the model favors variables which reflect building size (BuiltFAR, FacilFAR, BldgDepth) as well as approximations for previous sale prices (Last_Sale_Price and Last_Sale_Date). The classification model tends to favor spatial lag features, such as how many buildings were sold in the past year within 500 meters (Percent_Neighbors_Sold and Radius_Res_Units_Sold_In_Year)

as well as characteristics of the building function (Percent_Office, Percent_Storage, etc.).

5 Future Research and Conclusions

5.1 Future Research

This research has shown that the addition of spatial-lag features can increase the predictive accuracy of machine learning models above and beyond traditional spatial aggregation techniques. There are several areas that could be further explored regarding spatial lag features, some of which are mentioned below.

First, it became apparent in the research that generalization was a problem for some of the models, likely due to overfitting of the training data. This issue was corrected by employing different algorithms, however, further work could be done to create variable selection processes and/or hyperparameter tuning to prevent overfit.

Additionally, the spatial lag features seemd to perform best for certain boroughs and for residential building types. One possible explanation for this is that these types of assetts are considerably more numerous and homogenous. We hypothesize that using a 500 meter radius to build spatial lag features, a distance which was arbitrarily chosen, works best for this type of asset in these areas. Fotheringham (2015) used an "Adaptive Bandwidth" technique to adjust the spatial lag radius based on cross-validation with much success. The techniques presented in this paper could be expanded to use cross validation in a similar manner to assign the optimal spatial lag radius for each building type and location.

Finally, this research aimed to predict real estate transactions 1 year into the future. While this is a promising start, 1-year of lead time may not be sufficient to respond to growing gentirifaction challenges. In addition, modeling at the annual level could be improved to quartlery or monthly, given that the sales data contains date information down to the day. To make this system practial for combatting displacement, it may be helpful to predict at a more granular level and further into the future.

5.2 Conclusion

Gentification is largely beneficial to societies and communities, however, the downside should not be overlooked. Displacement causes Economic Exclusion, which over time can contribute to rising Income Inequality. Combatting displacement allows communities to benefit from gentrification without suffering the negative consequences. One way to practically combat displacement is to predict gentirification, which this paper has attempted to do.

Spatial lags, typically seen in geographically weighted regression, were employed successfully to enhance the predictive power of machine learning models. THe spatial lag models performed best for particular building types and geographies, however, we feel confident that the techique could be expanded to work equally as well for all buildings with some additional research. Regarding algorithms, Artificaly Neural Networks performed the best for predicting sale price, while GBM performed best for predicting sale occurance.

While this research is not intended to serve as a full early-warining system for gentrification and displacmenet, it is a step in that direction. More research is needed to help address the challenges faced by city planners and governments trying to help incumbent residents reap the benefits of local investments. Income inequallity is a complicated and grave issue, but new tools and tecnhiques to inform and prevent will help ensure equality of opportunity for all.

References

Almanie, R.; Lor, T.; Mirza. 2015. "Crime Prediction Based on Crime Types and Using Spatial and Temporal Criminal Hotspots." *International Journal of Data Mining & Knowledge Management Process (IJDKP)* 5 (4).

Antipov, Evgeny A., and Elena B. Pokryshevskaya. 2012. "Mass Appraisal of Residential Apartments: An Application of Random Forest for Valuation and a Cart-Based Approach for Model Diagnostics." *Expert Systems with Applications.*

Barry Bluestone & Chase Billingham, Stephanie Pollack &. 2010. "Maintaining Diversity in America's Transit-Rich Neighborhoods: Tools for Equitable Neighborhood Change." *New England Community Developments, Federal Reserve Bank of Boston*, 1–6.

Batty, Michael. 2013. "The New Science of Cities." *MIT Press.*

Breiman, L. 2001. "Random Forests." *Machine Learning* 45 (1): 5–32.

Breiman, Leo. 2001. "Random Forests." *Machine Learning* 45 (1). Springer: 5–32.

Chapple, Karen. 2009. "Mapping Susceptibility to Gentrification: The Early Warning Toolkit." *Berkeley, CA: Center for Community Innovation.*

Chapple, Miriam, Karen; Zuk. 2016. "Forewarned: The Use of Neighborhood Early Warning Systems for Gentrification and Displacement." *Cityscape: A Journal of Policy Development and Research* 18 (3).

Clay, Phillip L. 1979. *Neighborhood Renewal: Middle-Class Resettlement and Incumbent Upgrading in American Neighborhoods.* Lexington Books.

d'Amato, Tom, Maurizio; Kauko, ed. 2017. *Advances in Automated Valuation Modeling.* Springer International Publishing.

Dietzell, Nicole; Schäfers, Marian Alexander; Braun. 2014. "Sentiment-Based Commercial Real Estate Forecasting with Google Search Volume Data." *Journal of Property Investment & Finance,* 32 (6): 540–69.

Dreier, John; Swanstrom, Peter; Mollenkopf. 2004. *Place Matters: Metropolitics for the Twenty-First Century.* University Press of Kansas.

Eckert, J. K. 1990. *Property Appraisal and Assessment Administration.* Chicago, IL.: International Association of Assessing Officers.

Fotheringham, R; Yao, A.S.; Crespo. 2015. "Exploring, Modelling and Predicting Spatiotemporal Variations in House Prices." *The Annals of Regional Science* 54.

Friedman, Jerome H. 2001. "Greedy Function Approximation: A Gradient Boosting Machine." *The Annals of Statistics* 29 (5): 1189–1232.

Fu, Yanjie; et al. 2014. *Exploiting Geographic Dependencies for Real Estate Appraisal: A Mutual Perspective of Ranking and Clustering.* Proceedings of the 20th ACM SIGKDD international conference on Knowledge discovery; data mining.

Geltner, David, and Alex Van de Minne. 2017. "Do Different Price Points Exhibit Different Investment Risk and Return Commercial Real Estate." Real Estate Research Institute.

Guan, Jian, Donghui Shi, Jozef M. Zurada, and Alan S. Levitan. 2014. "Analyzing Massive Data Sets: An Adaptive Fuzzy Neural Approach for Prediction, with a Real Estate Illustration." *Journal of Organizational Computing and Electronic Commerce* 24 (1). Taylor & Francis: 94–112. https://doi.org/10.1080/10919392.2014.866505.

Hastie, Trevor, Robert Tibshirani, and Jerome Friedman. 2001. *The Elements of Statistical Learning.* Springer Series in Statistics. New York, NY, USA: Springer New York Inc.

Helbich, et al., Marco. 2013. "Boosting the Predictive Accuracy of Urban Hedonic House Price Models Through Airborne Laser Scanning." *Computers, Environment and Urban Systems* 39: 81–92.

Huang, Chong-Wei. 1996. "On the Complexity of Point-in-Polygon Algorithms." *Computers and Geosciences* 23.

Johnson, Ken, Justin Benefield, and Jonathan Wiley. 2007. "The Probability of Sale for Residential Real Estate." *Journal of Housing Research* 16 (2): 131–42. https://doi.org/10.5555/jhor.16.2.0234g75800h5k8x6.

Kontrimasa, Antanas, Vilius; Verikasb. 2011. "The Mass Appraisal of the Real Estate by Computational Intelligence." *Applied Soft Computing.*

Koschinsky, J. et al. 2012. "The Welfare Benefit of a Home's Location: An Empirical Comparison of Spatial and Non-Spatial Model Estimates." *Journal of Geographical Systems* 10109.

Lees, Tom; Wyly, Loretta; Slater. 2008. "Gentrification." *Growth and Change* 39 (3): 536–39. https://doi.org/10.1111/j.1468-2257.2008.00443.x.

Miller, J.; Aspinall, J.; Franklin. 2007. "Incorporating Spatial Dependence in Predictive Vegetation Models." *Ecological Modelling* 202 (3): 225–42.

Park, Jae Kwon, Byeonghwa; Bae. 2015. "Using Machine Learning Algorithms for Housing Price Prediction: The Case of Fairfax County, Virginia Housing Data." *Expert Systems with Applications* 42 (6): 2928–34.

Pivo, Gary, and Jeffrey D. Fisher. 2011. "The Walkability Premium in Commercial Real Estate Investments." *Real Estate Economics* 39 (2): 185–219. https://doi.org/10.1111/j.1540-6229.2010.00296.x.

Quintos, Carmela. 2013. "Estimating Latent Effects in Commercial Property Models." *Journal of Property Tax Assessment & Administration* 12 (2).

Rafiei, Hojjat, Mohammad Hossein; Adeli. 2016. "A Novel Machine Learning Model for Estimation of Sale Prices of Real Estate Units." *Journal of Construction Engineering and Management* 142 (2).

Reardon, Kendra, Sean F.; Bischoff. 2011. "Income Inequality and Income Segregation." *American Journal of Sociology.*

Ritter, Nancy. 2013. "Predicting Recidivism Risk: New Tool in Philadelphia Shows Great Promise." *National Institute of Justice Journal* 271.

Schernthanner H., Gonschorek J., Asche H. 2016. "Spatial Modeling and Geovisualization of Rental Prices for Real Estate Portals." *Computational Science and Its Applications* 9788.

Silverherz, J. D. 1936. "The Assessment of Real Property in the United States." *Albany:*

J.B. Lyon Co. Printers.

Smith, Neil. 1979. "Toward a Theory of Gentrification a Back to the City Movement by Capital, Not People." *Journal of the American Planning Association* 45 (4). Routledge: 538–48. https://doi.org/10.1080/01944367908977002.

Solomon Greene, Molly Scott, Rolf Pendall, and Serena Lei. 2016. "Open Cities: From Economic Exclusion to Urban Inclusion." *Urban Institue Brief.* Urban Institue Brief.

Turner, Margery Austin, and Christopher Snow. 2001. *Leading Indicators of Gentrification in D.c. Neighborhoods.*

Watson, Tara. 2009. "Inequality and the Measurement of Residential Segregation by Income in American Neighborhoods." *Review of Income and Wealth.*

Zuk, Miriam; et al. 2015. "Gentrification, Displacement and the Role of Public Investment: A Literature Review."